

In [7]: `from keras.datasets import imdb`

*# Load the data, keeping only 10,000 of the most frequently occurring words*  
`(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words = 10000)`

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>)  
 17464789/17464789 [=====] - 8s 0us/step

In [8]: *# Check the first review*  
`train_data[0]`

Out[8]: [1,  
 14,  
 22,  
 16,  
 43,  
 530,  
 973,  
 1622,  
 1385,  
 65,  
 458,  
 4468,  
 66,  
 3941,  
 4,  
 173,  
 36,  
 256,  
 5,  
 ...]

In [9]: *# Check the first Label*  
`train_labels[0]`

Out[9]: 1

In [10]: *# Since we restricted ourselves to the top 10000 frequent words, no word index should be out of range. We'll verify this below*

*# Here is a list of maximum indexes in every review --- we search the maximum index in each review*  
`print(type([max(sequence) for sequence in train_data]))`

*# Find the maximum of all max indexes*  
`max([max(sequence) for sequence in train_data])`

<class 'list'>

Out[10]: 9999

```
In [11]: # Let's quickly decode a review

# step 1: Load the dictionary mappings from word to integer index
word_index = imdb.get_word_index()

# step 2: reverse word index to map integer indexes to their respective words
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])

# Step 3: decode the review, mapping integer indices to words
#
# indices are off by 3 because 0, 1, and 2 are reserved indices for "padding", "Star
decoded_review = ' '.join([reverse_word_index.get(i-3, '?') for i in train_data[0]])

decoded_review
```

Downloading data from [https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb\\_word\\_index.json](https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json) ([https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb\\_word\\_index.json](https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json))

1641221/1641221 [=====] - 3s 2us/step

```
Out[11]: "? this film was just brilliant casting location scenery story direction everyone's
really suited the part they played and you could just imagine being there robert ? i
s an amazing actor and now the same being director ? father came from the same scott
ish island as myself so i loved the fact there was a real connection with this film
the witty remarks throughout the film were great it was just brilliant so much that
i bought the film as soon as it was released for ? and would recommend it to everyon
e to watch and the fly fishing was amazing really cried at the end it was so sad and
you know what they say if you cry at a film it must have been good and this definite
ly was also ? to the two little boy's that played the ? of norman and paul they were
just brilliant children are often left out of the ? list i think because the stars t
hat play them all grown up are such a big profile for the whole film but these child
ren are amazing and should be praised for what they have done don't you think the wh
ole story was so lovely because it was true and was someone's life after all that wa
s shared with us all"
```

```
In [12]: len(reverse_word_index)
```

```
Out[12]: 88584
```

```
In [13]: import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension)) # Creates an all zero matrix o
    for i,sequence in enumerate(sequences):
        results[i,sequence] = 1 # Sets specific indices of res
    return results

# Vectorize training Data
X_train = vectorize_sequences(train_data)

# Vectorize testing Data
X_test = vectorize_sequences(test_data)
```

```
In [14]: X_train[0]
```

```
Out[14]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
In [15]: X_train.shape
```

```
Out[15]: (25000, 10000)
```

```
In [16]: y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

```
In [17]: from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
In [19]: from keras import optimizers
from keras import losses
from keras import metrics

model.compile(optimizer=optimizers.RMSprop(learning_rate=0.001),
              loss = losses.binary_crossentropy,
              metrics = [metrics.binary_accuracy])
```

```
In [20]: # Input for Validation
X_val = X_train[:10000]
partial_X_train = X_train[10000:]

# Labels for validation
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

```
In [21]: history = model.fit(partial_X_train,
                             partial_y_train,
                             epochs=20,
                             batch_size=512,
                             validation_data=(X_val, y_val))
```

```
Epoch 1/20
30/30 [=====] - 2s 45ms/step - loss: 0.5339 - binary_accuracy: 0.7831 - val_loss: 0.4112 - val_binary_accuracy: 0.8661
Epoch 2/20
30/30 [=====] - 1s 29ms/step - loss: 0.3329 - binary_accuracy: 0.9001 - val_loss: 0.3272 - val_binary_accuracy: 0.8819
Epoch 3/20
30/30 [=====] - 1s 23ms/step - loss: 0.2448 - binary_accuracy: 0.9233 - val_loss: 0.2900 - val_binary_accuracy: 0.8872
Epoch 4/20
30/30 [=====] - 1s 21ms/step - loss: 0.1912 - binary_accuracy: 0.9413 - val_loss: 0.2820 - val_binary_accuracy: 0.8853
Epoch 5/20
30/30 [=====] - 1s 23ms/step - loss: 0.1540 - binary_accuracy: 0.9533 - val_loss: 0.2762 - val_binary_accuracy: 0.8913
Epoch 6/20
30/30 [=====] - 1s 24ms/step - loss: 0.1265 - binary_accuracy: 0.9615 - val_loss: 0.2922 - val_binary_accuracy: 0.8862
Epoch 7/20
30/30 [=====] - 1s 21ms/step - loss: 0.1034 - binary_accuracy: 0.9717 - val_loss: 0.2992 - val_binary_accuracy: 0.8857
Epoch 8/20
30/30 [=====] - 1s 21ms/step - loss: 0.0849 - binary_accuracy: 0.9773 - val_loss: 0.3312 - val_binary_accuracy: 0.8769
Epoch 9/20
30/30 [=====] - 1s 24ms/step - loss: 0.0698 - binary_accuracy: 0.9827 - val_loss: 0.3386 - val_binary_accuracy: 0.8811
Epoch 10/20
30/30 [=====] - 1s 22ms/step - loss: 0.0561 - binary_accuracy: 0.9871 - val_loss: 0.3628 - val_binary_accuracy: 0.8815
Epoch 11/20
30/30 [=====] - 1s 23ms/step - loss: 0.0444 - binary_accuracy: 0.9912 - val_loss: 0.4008 - val_binary_accuracy: 0.8783
Epoch 12/20
30/30 [=====] - 1s 21ms/step - loss: 0.0372 - binary_accuracy: 0.9927 - val_loss: 0.4163 - val_binary_accuracy: 0.8769
Epoch 13/20
30/30 [=====] - 1s 22ms/step - loss: 0.0265 - binary_accuracy: 0.9959 - val_loss: 0.4471 - val_binary_accuracy: 0.8753
Epoch 14/20
30/30 [=====] - 1s 21ms/step - loss: 0.0220 - binary_accuracy: 0.9970 - val_loss: 0.5078 - val_binary_accuracy: 0.8718
Epoch 15/20
30/30 [=====] - 1s 21ms/step - loss: 0.0183 - binary_accuracy: 0.9965 - val_loss: 0.5067 - val_binary_accuracy: 0.8727
Epoch 16/20
30/30 [=====] - 1s 22ms/step - loss: 0.0118 - binary_accuracy: 0.9987 - val_loss: 0.5413 - val_binary_accuracy: 0.8724
Epoch 17/20
30/30 [=====] - 1s 21ms/step - loss: 0.0103 - binary_accuracy: 0.9987 - val_loss: 0.5760 - val_binary_accuracy: 0.8708
Epoch 18/20
30/30 [=====] - 1s 23ms/step - loss: 0.0053 - binary_accuracy: 0.9996 - val_loss: 0.6348 - val_binary_accuracy: 0.8646
Epoch 19/20
30/30 [=====] - 1s 25ms/step - loss: 0.0055 - binary_accuracy: 0.9995 - val_loss: 0.6613 - val_binary_accuracy: 0.8641
Epoch 20/20
30/30 [=====] - 1s 24ms/step - loss: 0.0027 - binary_accuracy: 0.9999 - val_loss: 0.6845 - val_binary_accuracy: 0.8675
```

```
In [22]: history_dict = history.history
history_dict.keys()
```

```
Out[22]: dict_keys(['loss', 'binary_accuracy', 'val_loss', 'val_binary_accuracy'])
```

```
In [23]: import matplotlib.pyplot as plt
%matplotlib inline
```


```
In [24]: # Plotting Losses
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, 'bo', label="Training Loss")
plt.plot(epochs, val_loss_values, 'b', label="Validation Loss")

plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss Value')
plt.legend()

plt.show()
```



In [25]: # Training and Validation Accuracy


```
acc_values = history_dict['binary_accuracy']
val_acc_values = history_dict['val_binary_accuracy']

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, acc_values, 'ro', label="Training Accuracy")
plt.plot(epochs, val_acc_values, 'r', label="Validation Accuracy")

plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



```
In [26]: model.fit(partial_X_train,
                    partial_y_train,
                    epochs=3,
                    batch_size=512,
                    validation_data=(X_val, y_val))
```

```
Epoch 1/3
30/30 [=====] - 1s 33ms/step - loss: 0.0021 - binary_accuracy: 0.9999 - val_loss: 0.7997 - val_binary_accuracy: 0.8623
Epoch 2/3
30/30 [=====] - 1s 21ms/step - loss: 0.0021 - binary_accuracy: 0.9998 - val_loss: 0.7724 - val_binary_accuracy: 0.8667
Epoch 3/3
30/30 [=====] - 1s 25ms/step - loss: 8.9476e-04 - binary_accuracy: 1.0000 - val_loss: 0.8192 - val_binary_accuracy: 0.8670
```

Out[26]: <keras.callbacks.History at 0x1f286022eb0>

```
In [35]: # Making Predictions for testing data
np.set_printoptions(suppress=True)
result = model.predict(X_test)
```

782/782 [=====] - 2s 2ms/step

```
In [36]: result
```

```
Out[36]: array([[0.00977155],
                [1.         ],
                [0.9957733 ],
                ...,
                [0.00096024],
                [0.00696085],
                [0.8679086 ]], dtype=float32)
```

```
In [37]: y_pred = np.zeros(len(result))
for i, score in enumerate(result):
    y_pred[i] = 1 if score > 0.5 else 0
```

```
In [38]: from sklearn.metrics import mean_absolute_error
mae = mean_absolute_error(y_pred, y_test)
```

```
In [39]: mae
```

```
Out[39]: 0.1472
```